
SaGess: Sampling Graph Denoising Diffusion Model for Scalable Graph Generation

Stratis Limnios

The Alan Turing Institute
London, UK
slimnios@turing.ac.uk

Praveen Selvaraj

The Alan Turing Institute
London, UK
pselvaraj@turing.ac.uk

Mihai Cucuringu

Department of Statistics & Mathematical Institute
University of Oxford &
The Alan Turing Institute, London, UK
mihai.cucuringu@stats.ox.ac.uk

Carsten Maple

University of Warwick
Coventry, UK
cm@warwick.ac.uk

Gesine Reinert

Department of Statistics
University of Oxford &
The Alan Turing Institute, London, UK
reinert@stats.ox.ac.uk

Andrew Elliott

Department of Mathematics and Statistics
University of Glasgow &
The Alan Turing Institute, London, UK
Andrew.Elliott@glasgow.ac.uk

Abstract

Over recent years, denoising diffusion generative models have come to be considered as state-of-the-art methods for synthetic data generation, especially in the case of generating images. These approaches have also proved successful in other applications such as tabular and graph data generation. However, **due to computational complexity, to this date, the application of these techniques to graph data has been restricted to small graphs, such as those used in molecular modeling.** In this paper, we propose SAGESS, a discrete denoising diffusion approach, which is able to generate large real-world networks by augmenting a diffusion model (DIGRESS) with a generalized divide-and-conquer framework. The algorithm is capable of generating larger graphs by sampling a covering of subgraphs of the initial graph in order to train DIGRESS. SAGESS then constructs a synthetic graph using the subgraphs that have been generated by DIGRESS. We evaluate the quality of the synthetic data sets against several competitor methods by comparing graph statistics between the original and synthetic samples, as well as evaluating the utility of the synthetic data set produced by using it to train a task-driven model, namely link prediction. In our experiments, SAGESS outperforms most of the one-shot state-of-the-art graph generating methods by a significant factor, both on the graph metrics and on the link prediction task.

1 Introduction

Synthetic data are key to many methods in machine learning and statistics, and synthetic data generation has sparked a significant amount of attention in recent years. Tools such as Dall-E for synthetic image generation, and ChatGPT have triggered the curiosity of the general public; the majority of the machine learning community are also diverting interest and resources towards generative algorithms. Beyond generating appealing synthetic images or asking ChatGPT to write

prose, many real world applications benefit greatly from synthetic data generation, for tasks including data augmentation in the training of classifiers/regressors[22], privacy protection of sensitive data[3] or removing bias from data sets [20]. Synthetic graph generation for modeling social interactions, generating new chemical compounds, or forecasting transactions are capital tasks that require efficient methods.

Many complex data sets can be represented as networks, and hence synthetic graph generators are of particular interest. Often these networks are viewed as realizations of a random process. The design of generative models for random graphs has a rich history, coming from probabilistic and structural assumptions with traditional methods such as early work in Erdős-Rényigraphs [5], stochastic block models [10], exponential random graphs [16] or the Barabasi-Albert model [1]. However, these methods often oversimplify the underlying complex structure of the graphs and are often not able to capture the distributions arising from real-world scenarios. Thus, more recently there has been an increasing interest in the community in designing deep models for synthetic graph generation, which allows for more flexible algorithms that are able to capture the intricacies of real networks with complex dependencies between the edges. Many approaches have been proposed for graph generation including autoregressive approaches which generate nodes and edges step by step. Such methods include GraphRNN [27] and GRAN [12], which improves modeling of long-term dependencies using a graph-based attention mechanism. Additional approaches include autoencoder based approaches, such as GraphVAE [18], adversarial approaches such as [4]. There are many more methods; detailed reviews are found for example in [6, 7].

Moreover, after the outstanding performances of Denoising Diffusion Models [9, 19] on image generation, various lines of work have applied such models to graph learning, mostly for molecular generation such as GeoDiff [23] or chemical compound design [17]. More versatile models have been introduced, such as DIGRESS [21], a discrete denoising diffusion graph model that is able to generate very high quality graphs with node and edge attributes. One of the main assets of denoising diffusion models is not having to rely on adversarial training, but they still need a large data set to train. The implication of this is that models such as DIGRESS do not adapt well when the task is to generate one large graph from a single sample.

Formally, creating a synthetic graph generator is equivalent, either explicitly or implicitly, to estimating/sampling from a probability distribution over space of possibly directed, possibly weighted/attributed graphs. When many independent realizations from the unknown distribution are available, then standard estimation methods are often successful. Yet, many real world graph data sets consist of only a single graph, either because it is expensive to measure, or because of the nature of the data, for example a global social network. Thus, there is interest in the more difficult task of approximating the underlying probability distribution only viewing a single sample. Classical models often navigate this task by making strong assumptions, such as Erdős-Rényi, where we assume the edges are i.i.d.. Recent deep learning methods often approach this challenge by the structured model of a GAE/GVAE or by learning the distribution over subsamples of nodes and edges; an example is ¹ e.g. NetGAN, an adversarial approach which relies on random walks [4].

In particular, while DIGRESS is a powerful approach for generating small graphs, its requirement to have multiple training samples as input prevents its application to the situation when the input is just a single network. Building on DIGRESS, in this work we propose a denoising-based diffusion model that can operate in the single large graph data set case. To achieve this, we leverage the strengths and the quality of the synthetic graph data produced by DiGress, in order to design a denoising diffusion based model. Along the way, we alleviate two of its weaknesses in this context, not only its requirement for a large number of training samples, which are not available in this case, but also its poor scalability to larger graphs.

We achieve this by leveraging a divide-and-conquer scheme, which allows us to extend this method to larger single graphs.

Our main contributions can be summarized as follows.

- We introduce graph subsampling methods, to break large graphs into a trainable data set.
- We propose SAGESS (SAmpling Graph dEnoiSing DiffuSion model), a pipeline that allows us to employ DIGRESS as a sample generating base for larger graphs.

¹In effect, this assumes that the underlying process that generated the network is related to these samples.

- We propose a task-driven evaluation by training link prediction GVAE on synthetic data and testing on real.

The paper is structured as follows. In Section 2 we set up the notations for the rest of the paper, as well as the DIGRESS model that is at the center of our framework, while also elaborating on its limitations. In Section 3, we state the problem we aim to address, and then present our solution framework SAGESS, which unfolds in two sections: first, the sampling methods to obtain a training data set, and second, the reconstruction pipeline of the synthetic graph. Next, in sections 4 and 5, we respectively present our experimental setup and results. Finally, we discuss potential limitations and future work in Section 6.

2 Graph Diffusion Model and Scalability

Notations and definitions presented in this section will be used throughout the rest of the paper. We provide the essential framework to support our work, and elaborate on the problems we address with the introduction of the SAGESS framework.

Graphs in this paper are denoted as $G = (V, E)$ where $|V| = n$ is the set of nodes and $|E| = e$ is the set of edges. Each node is of one of a types, and each edge is of one of b types. We associate with G the matrices $\mathbf{X} \in \mathbf{R}^{n \times a}$ where $\mathbf{X}_{i,:}$ is the one-hot encoding for the feature of node i , and $\mathbf{E} \in \mathbf{R}^{n \times n \times b}$ where $\mathbf{E}_{i,j,:}$ is the one-hot encoding for the feature of the edge between nodes i and j . We denote by $P_k(V)$ the set of all k -point subsets of V and by $S_k(G)$ the set of all possible subgraphs of G of size k . The subgraph of G induced by the nodes in S is denoted by $G[S]$; we also refer to such graphs as *node induced subgraphs*.

2.1 Graph Generation using Discrete Denoising Diffusion Model: DIGRESS

Here we introduce the key aspects of Denoising Diffusion Probabilistic Models (DDPM) with particular emphasis on application to graphs. We also point out limitations of state-of-the art graph generation using discrete space diffusion models.

DIGRESS [21] is currently one of the most efficient tools in graph generation. Taking as input a data set of a variety of graphs, it learns a denoising process in discrete space and is able to mimic the input graphs with remarkable precision. The key aspect of this method, and more generally graph generation methods based on diffusion models [8, 21], relies on a discrete space noise scheduling. Indeed, instead of the standard Gaussian noising and denoising procedure, these frameworks propose to add iteratively discrete noise via edge and node addition and deletion at random.

Traditionally, these models are based on a forward and a reverse Markov process. Indeed a forward process denoted by $q(A^{1:T}|A_0) = \prod_{t=1}^T q(A^t|A^{t-1})$ generates an increasingly noisier samples from the candidate A^0 to white noise A^T , A^t here being adjacency matrices and T is a hyperparameter; in DIGRESS the default is 500. We then learn the reverse process $p_\theta(A^{1:T}) = p(A^T) \prod_{t=1}^T q(A^{t-1}|A^t)$ that aims to denoise the latent adjacency matrices A^t to produce the synthetic samples.

The DIGRESS model takes as state space the set of node types and of edge types. DIGRESS then defines the transition probabilities from one state to another for nodes and edges through the noise matrices $[Q_X]_{i,j}^t = q(x^t = j|x^{t-1} = i)$ and $[Q_E]_{i,j}^t = q(e^t = j|e^{t-1} = i)$ where Q is chosen so that the Markov chain converges to the relative type frequencies in the initial population. Then, to get a noisy sample $G^t = (\mathbf{X}^t, \mathbf{E}^t)$ each node and each edge type is sampled from the categorical distribution $q(G^t|G) = (\mathbf{X}\bar{\mathbf{Q}}_X^t, \mathbf{E}\bar{\mathbf{Q}}_E^t)$ with $\bar{\mathbf{Q}}_X^t = \mathbf{Q}_X^1 \dots \mathbf{Q}_X^t$ and $\bar{\mathbf{Q}}_E^t = \mathbf{Q}_E^1 \dots \mathbf{Q}_E^t$.

The denoising component of the DIGRESS model is a denoising neural network ϕ_θ parametrized by θ . It is trained by optimizing the cross-entropy loss between the predicted probabilities $\hat{p} = (\hat{p}^X, \hat{p}^E)$ and the true graph G , and is defined by

$$l(\hat{p}^G, G) = \sum_{1 \leq i \leq n} \text{cross-entropy}(x_i, \hat{p}_i^X) + \lambda \sum_{1 \leq i \leq n} \text{cross-entropy}(e_{ij}, \hat{p}_{ij}^E), \quad (1)$$

where $\lambda \in \mathbb{R}^+$ balances the importance between nodes and edges.

Simply put, DIGRESS is a way to obtain a classification for each node and edge.

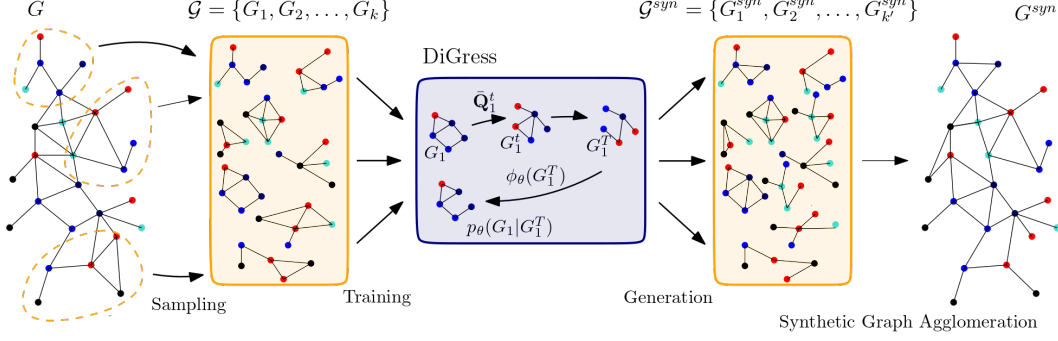


Figure 1: SAGESS Diagram

Finally, after training the network, it can be used to generate synthetic graphs via the estimation of the reverse diffusion iteration $p_\theta(G^t|G)$ using \hat{p}^G , as follows. We put $p_\theta(x_i^{t-1}|x_i = x, G^t) = q(x_i^{t-1}|x_i = x, x_i^t)\mathbf{1}(q(x_i^t|x_i = x) > 0)$ and $p_\theta(e_{ij}^{t-1}|e_{ij} = e, G^t) = q(e_{ij}^{t-1}|e_{ij} = e, e_{ij}^t)\mathbf{1}(q(e_{ij}^t|e_{ij} = e) > 0)$, then set

$$p_\theta(x_i^{t-1}|G^t) = \sum_x p_\theta(x_i^{t-1}|x_i = x, G^t)\hat{p}_i^X(x), \quad p_\theta(e_{ij}^{t-1}|G^t) = \sum_x p_\theta(e_{ij}^{t-1}|e_{ij} = e, G^t)\hat{p}_{ij}^E(e),$$

and finally obtain $p_\theta(G^{t-1}|G^t) = \prod_{i=1}^n p_\theta(x_i^{t-1}|G^t) \prod_{1 \leq i, j \leq n} p_\theta(e_{ij}^{t-1}|G^t)$. The sampled G^{t-1} then serves as input of the denoising network at the next time step. This derivation shows the complexity of the calculations; when the numbers a and b of types are small, when the network is small, and when T is small, they are feasible, but large networks with many different types can pose a challenge. More details are in the SI.

Overall, DIGRESS can perform extremely well to generate small graphs with attributes from a finite set of attributes. Unfortunately, the complexity of DIGRESS and the way the model is set up in a way that prevents us from using it to sample one large graph. The complexity lies in the dimensions of the matrices \mathbf{X} , $\bar{\mathbf{Q}}_X^t$, \mathbf{E} and $\bar{\mathbf{Q}}_E^t$ as at least the edge matrices scale in n^2 times the number of attributes, it is indeed specified in the [21] that DIGRESS has complexity $\Theta(n^2)$ per layer, due to the attention scores and the predictions for each edge. Hence, not only we need to store in memory these matrices for every diffusion step, but we also need to learn T dense matrices Q_E to define the Markov transitions. This is fairly reasonable for small graphs, but starts to become computationally difficult to approach when the number of nodes in the graph increases. Then, supposing we are able to run DIGRESS on large graphs, we cannot train it on one single graph. DIGRESS, much like other deep-learning algorithms, thrives when having access to a large training set. Altogether, this provides the motivation for proposing our framework based on sampling from large graphs.

3 Sampling Graph Denoising Diffusion Model Generator (SAGESS)

This paper addresses the task to generate a single graph from a unique observation. We do not have a large training set of graphs \mathcal{G} available to train any type of diffusion model. Our solution, which we denote as SAGESS and which is visualized in Figure 1, can be summarized as follows.

1. First, we produce a graph data set \mathcal{G} issued from the initial graph G .
2. Next, we employ DiGress to create samples using the data set \mathcal{G} as training data.
3. Finally, we rebuild a graph from the trained diffusion model in a systematic fashion, paying particular care to match the node ids in the samples.

3.1 Graph Sampling and Covering

The first step is to construct a representative data set to train DiGress. We propose three sampling schemes to capture different structural properties of the graph. Each method addresses one or more

properties of the graph, from local to more global ones. Depending on the real world graph we apply our model to, some sampling schemes might be more effective than others; for instance, ego networks might benefit more from locality in the sampling, whereas communication data sets, for instance, might benefit more from more global features such as random walks.

Uniform Node Sampling (Unif): We sample a set of node induced subgraphs by selecting k sized subsets of nodes S_k uniformly at random from $P_k(V)$. We note that this sampling scheme is invariant under permutation of the nodes. Based on results from [14], if p denotes the probability of selecting a node, then $\Theta(p^{-2} \log n \log(1/\delta))$ uniformly node induced subgraphs suffice for reconstruction with probability at least $1 - \delta$. In our setting, $p = k/n$. To give a broad sense to the number of subgraphs we need to sample, for a graph with $n = 1,000$ nodes, suppose we select uniformly at random $2 \log(n)$ nodes. We would then need $\Theta(4 \log^3(n) \log(1/\delta))$ subgraphs of size 20 to reconstruct G . Setting the probability $1 - \delta = 0.95$ results in a data set size of the order of 10,000 graphs. Here we use 10,000 for simplicity; other values are explored in the SI.

Random Walk Node Induced Subgraph (RW): This second method is based on sampling from subgraphs induced by random walks starting from every node in the graph G . The idea is straightforward: for each node v_1 of the n nodes in V we generate a k step random walk $w^k = \{v_1, \dots, v_k\}$ starting from v_1 . Then we obtain the node induced subgraph from those walks $G[w_k]$. We repeat this construction d times. This gives us a data set $\mathcal{G} = \{G[w_1^k], \dots, G[w_{d \times n}^k]\}$, where d is a hyper-parameter controlling the number of sampled graphs per node, see supplementary material for additional discussion.

2-hop Neighborhood Sampling (Ego): The last sampling method is based on sampling from every node a 2-hop neighborhood with random node deletion. As a 2-hop neighborhood can cover most of the graph and we aim for small samples to be able to train DIGRESS, with ideally more than one subgraph including each node, we fix to k the maximum size of the subgraph we generate from the 2-hop neighborhood of a node. To do so, if the neighborhood is larger than k , we delete (the integer part of) half of nodes from the neighborhood uniformly at random. We also ensure to keep, at every step, the largest connected component until there are less than k nodes left; we denote this resulting node set by

$\mathcal{N}^k(v)$. Again, we obtain its node induced subgraph $G[\mathcal{N}^k(v)]$. For k small enough, there is randomness in our modified 2-hop neighborhoods. We generate d subgraphs per node, with d as in the previous paragraph, hence producing a data set $\mathcal{G} := \{G[\mathcal{N}_1^k(v_1)], \dots, G[\mathcal{N}_d^k(v_1)], \dots, G[\mathcal{N}_1^k(v_n)], \dots, G[\mathcal{N}_d^k(v_n)]\}$. Theoretical guarantees for this subsampling scheme are available for example in [2].

Algorithm 1 Ego Sampling

```

1: for  $i \in \{1, \dots, n\}$  do
2:   for  $j \in \{1, \dots, d\}$  do
3:      $G_i = \mathcal{N}(v_i)$ 
4:     while  $|V_i| > k$  do
5:       choose uniformly  $\lfloor \frac{|V_i|}{2} \rfloor$  nodes from  $V_i$ ; call this set  $S_i$ 
6:       delete the nodes in  $S_i$  from  $G_i$ 
7:       retain the largest connected component of  $G_i$ 
8:     end while
9:   end for
10: end for
```

Each of these methods provides a graph training set for DIGRESS which is obtained from one single graph. Moreover, following [14], we produce enough samples to have a good chance to cover each node and edge in the uniform sampling scheme, with the heuristic that this should also be adequate for the denser sampling schemes. We also shuffle the training data set \mathcal{G} before feeding it to DIGRESS. Next, we detail how SAGESS reconstructs the initial graph from the trained model.

3.2 Node Labeling and Reconstruction

Our SAGESS model first samples subgraphs to train DIGRESS, but DIGRESS will only be able to generate small graphs from the training set \mathcal{G} it learns from. This would lead to a non-identifiable set of small graphs, that one would not know how to merge into one single graph.

This is where the ability of DIGRESS to handle node attributes is important. We set the initial node ids as node features on the subgraphs to enable identification after generating synthetic samples. In detail, to every graph $G_i^k \in \mathcal{G}$ we associate a feature matrix $\mathbf{X}_i \in \mathbb{R}^{k \times n}$, where k is the number of nodes in G_i^k and n the number of nodes in G , so that \mathbf{X}_i is a one-hot encoding of the node id in the initial graph. This enables the model to learn the local graph structure as we identify the nodes with the attribute. We then generate enough small graph samples with DIGRESS and agglomerate the new edges and nodes uniquely until we match the initial graph G 's edge count. The agglomeration step unfolds as follows. We start with the first generated graph. Then as long as the number of edges, $|E|$, in the original graph is not reached, we generate a new synthetic graph G_{syn}^k with DIGRESS and take the graph union with the current graph. The sampling process stop when $|E_{syn}| > |E|$; if G_{syn}^l is the last graph sampled, we add all the new edges to G_{syn} regardless of whether the new edge count exceeds $|E_{syn}|$. As the generated subgraphs tend to be small, the overshoot tends to be small also.

Algorithm 2 SAGESS

- 1: \mathcal{G} : Sample $n \times d$ samples using Unif/Ego/RW from G
 - 2: Train DIGRESS with \mathcal{G}
 - 3: **while** $|E_{syn}| < |E|$ **do**
 - 4: generate G_{syn}^k
 - 5: add new unique edges from (G_{syn}^k) to G_{syn}
 - 6: **end while**
-

As it builds on DIGRESS, SAGESS comes with theoretical guarantees inherited from DIGRESS, in particular the permutation equivariant architecture and a permutation invariant loss. Exchangeability of the generated distributions then follows under the exchangeable SAGESS sampling schemes, permuting the order of the samples in the input to destroy any potential sequential dependence. These properties ensure that SAGESS can learn efficiently from the data.

4 Experimental Evaluation

In this section we will be evaluating our framework against state-of-the-art graph generation methods on four real world data sets and one synthetic data set. First, we will compare graph statistics to evaluate the quality of the generated graphs. Then we will train a Variational Graph Auto-Encoder (GVAE) to evaluate the utility of the synthetic data generated on a link prediction task.

4.1 Benchmark models

To benchmark our approach, we will compare against several competitor methods which construct synthetic data based on a single sample. To make a meaningful comparison, we have selected methods from multiple different generation approaches, from classical approaches, through adversarial approaches, and approaches based on a low-rank approximation. They are as follows:

DCSBM [11] A classic approach from the network science literature, the so-called degree corrected stochastic block model² assumes a classic stochastic block model, where the nodes are divided into blocks, and the probability of a connection between a pair of nodes is a function of their block membership and additionally their degree.

NetGAN [4] An adversarial approach leveraging the GAN framework, to learn the distribution of biased random walks which are then combined via a transition matrix into a sampled graph. Edges in the sampled graph are sampled (mostly) uniformly at random.³

CELL [15] A modification of NetGAN which replaces the GAN formulation with a formulation

²Implementation from <https://github.com/microsoft/graspologic>

³Implementation from <https://github.com/danielzuegner/netgan>

based on a low-rank approximation, see paper for full discussion. In this formulation, sampled graph edges are again sampled in an edge independent manner.⁴

4.2 Data sets

We evaluate our method on four real world data sets from the torch geometric package ⁵ and one synthetic data set:

EuCore: This data set is an e-mail communication network of a large European research institution, from [26]. Nodes indicate members of the institution, and an edge between a pair of members indicates that they exchanged at least one email. This graph has 1005 nodes and 16,706 edges.

Cora: Cora is a citation data set from [25]. Every node is an article and an edge links two nodes if one cites the other. It consists of a directed graph with 2,708 nodes and 10,556 edges.

Wiki: A data set from wikipedia pages from [24] with 2,405 nodes and 12,761 undirected edges.

Facebook: This data set consists of friend lists from Facebook published in [13]. This data set initially contains 10 graphs, but we will only be using the largest one (second graph) which has 1,045 nodes and 27,755 undirected edges.

SBM: This is a standard Stochastic Block Model (SBM) generated graph, with 4 blocks of sizes 400. We set the inner-cluster density to 0.15 and the across-cluster density to 0.01.

4.3 Experiments

We evaluate our framework using three different experimental settings. In our first experiment, we evaluate how close the structural properties of the generated graphs from each of the methods are to the original graph. For evaluation we choose the set of standard metrics from Ref. [15] such as number of nodes (ignoring isolated nodes), clustering coefficient, assortativity or triangle count. As some generation methods fix the number of nodes and others fix the number of edges, we report these numbers also but do not assess the methods on them.

To further explore our approach, we additionally compare the effectiveness of each of our sampling schemes using the same approach. We produce and compare synthetic graphs for all data sets for the Random Walk (RW), and 2-hop modified neighborhood (Ego) sampling schemes. While for the Uniform (Unif) sampling scheme, we compare on EuCore, Facebook and SBM, noting that this sampling scheme struggles on the remaining data sets, due to their sparsity.

In essence, as the induced subgraphs in the training set from the uniform sampling tend to have very few edges (less than the nodes of the subgraph), our model will struggle to find the right amount of edges to match the edge count of the initial graph.

The second experiment consists of a utility test. Indeed, we want to evaluate the usefulness of the synthetic data set produced. Hence, we propose to train a Variational Graph Auto-Encoder for link prediction, the GVAE is trained on the generated synthetic data. Once trained, we obtain a latent space from the synthetic data set on the set of nodes belonging to the real data set. The evaluation is the following: we compute the probability of each possible edge using the GVAE encodings and evaluate them on the 90% of the edges of the real graph. This can also be interpreted as a test on how efficient the synthetic data set can be, when employed for a data augmentation task.

Finally, we also explore the ability of our method to produce meaningful smaller graph samples. Indeed, it might not be needed to generate a graph of the same size as the initial graph G . Our goal is to demonstrate that our model can also generate smaller scale synthetic graphs in terms of edges that are structurally similar to the initial graph. We generate graphs of sizes ranging from 10% to 100% of the number of edges in the initial graph, and observe the computed metrics convergences to the right values in Figure 2. All experimental hyperparameters, as well as extended experiments can be found in the supplementary.

⁴Implementation from <https://github.com/hheidrich/CELL>

⁵<https://pytorch-geometric.readthedocs.io/en/latest/modules/datasets.html>

Table 1: Comparing network statistics of graphs generated from each synthetic data method on five different graph data sets. This table includes the graph statistics of Table 2 in Ref. [15], from a similar experiment, with CPL denoting the characteristic path length; we also report the number of nodes and the number of edges. The text [ABC](#) in blue indicates the closest statistic to the real data, whereas the green text [ABC](#) indicates the second closest.

	Method	num nodes	num edges	num triangles	num squares	max deg	cluster coef.	assort.	power law exp	CPL
EuCore	Real	1005	16,706	105,461	4,939,311	346	0.39935	-0.01099	1.3621	2.58693
	SAGESS-Uni	939	16,716	114,900	6,280,664	287	0.34024	-0.06321	1.35697	2.49139
	SAGESS-RW	878	16,709	131,429	6,995,335	351	0.42338	-0.04234	1.35192	2.50512
	SAGESS-Ego	867	16,707	114,593	5,257,738	342	0.39000	-0.023969	1.3296	2.43646
	NetGAN	986	16,064	62,278	2,505,330	279	0.25569	-0.06196	1.34179	2.48189
	CELL	1005	16,064	74,251	3,336,294	273	0.29808	-0.07655	1.36279	2.56782
	DCSBM	951	15,906	75,743	3,699,308	305	0.19673	-0.010515	1.35087	2.47134
Cora	Real	2708	10,556	1,630	4,664	168	0.24067	-0.06587	1.93230	6.31031
	SAGESS-RW	2548	10,557	1,806	14,952	220	0.25728	-0.05969	1.84290	5.43069
	SAGESS-Ego	2540	10,562	1,696	7,923	171	0.23911	-0.06868	1.92240	5.77869
	NetGAN	2485	10,138	932	2,394	128	0.15687	-0.07384	1.86148	5.86039
	CELL	2708	10,556	521	1,295	97	0.07930	-0.08226	1.88642	6.03067
	DCSBM	2621	10,097	2,380	2,809	237	0.06991	-0.01709	1.85581	4.47947
Wiki	Real	2405	12,761	23,817	407,302	263	0.37581	-0.07875	1.54227	3.65161
	SAGESS-RW	2348	12,763	26,296	466,120	280	0.38389	-0.09282	1.54076	3.53405
	SAGESS-Ego	2275	12,763	26,897	506,620	312	0.37530	-0.10494	1.55261	3.47834
	NetGAN	2405	11,596	10,701	104,243	241	0.20179	-0.10076	1.52944	3.60293
	CELL	2357	11,592	10,136	121,014	258	0.22965	-0.10769	1.54027	3.62619
	DCSBM	2251	11,595	9,439	167,655	263	0.08492	-0.01821	1.54596	3.45205
Facebook	Real	1045	27,755	446,846	34,098,662	1044	0.57579	-0.02543	1.28698	1.94911
	SAGESS-Uni	1043	27,758	429,428	35,261,545	999	0.52098	-0.01607	1.29003	2.00800
	SAGESS-RW	1009	27,764	490,844	43,006,252	1001	0.56138	-0.02266	1.29398	1.96014
	SAGESS-Ego	1005	27,761	515,928	45,421,130	295	0.43074	0.34074	1.29381	2.65926
	NetGAN	1045	27,755	262,574	15,635,262	849	0.39773	-0.01821	1.27429	2.13730
	CELL	1045	27,755	250,968	14,855,676	474	0.30854	0.12788	1.27490	2.38650
	DCSBM	1041	27,092	339,448	26,714,948	733	0.37549	0.07125	1.28845	2.33021
SBM	Real	1600	73,312	344,574	20,955,308	155	0.15418	-0.00188	3.58894	2.08276
	SAGESS-Uni	1600	73,313	342,639	21,054,221	164	0.14830	0.03572	2.16732	2.04884
	SAGESS-RW	1600	73,367	490,663	34,384,218	207	0.20062	0.22337	1.70495	2.13963
	SAGESS-Ego	1600	73,326	366,162	22,979,792	174	0.15622	0.07177	1.89330	2.06118
	NetGAN	1600	73,312	367,143	22,775,320	144	0.16054	0.19039	2.39375	2.12747
	CELL	1600	73,312	341,351	20,783,575	139	0.15186	-0.00156	2.72805	2.08060
	DCSBM	1600	73,357	353,934	21,794,585	130	0.15671	-0.00343	2.79204	2.08274

5 Results

Throughout our experiments, we have shown that SAGESS is a very powerful tool which can create high quality synthetic graphs. This is demonstrated in our first experiment, where the statistics on the synthetic graphs (Table 1) we produce, perform well in comparison to the benchmark models, across all the benchmark data sets. In particular, the SAGESS-RW method performs well in the first four data sets, although struggles to some extent in the synthetic SBM data set. It is worth noticing as well that we tend to generate local structures, such as triangles, easier than the other baselines. This is due to the fact that DIGRESS is very efficient at sampling motifs, and the subgraphs used for training contain many of them due to the nature of the sampling, which samples subgraphs with interdependencies between the edges. It is important to point out that even if we train on local substructures via the subgraphs, we manage to accurately obtain global metrics like characteristic path length or clustering coefficient.

In our second experiment, we present the results in Table 2 where we exhibit the utility of our method that overall performs very consistently compared to the other methods. We also notice that while we are close to other methods in terms of statistics on some data sets, the utility performance on the link prediction is somewhat more diverse. For example, while the DCSBM method performs well in terms of graph statistics on the synthetic SBM data set (see Table 1), we still outperform it by a significant factor in the link prediction task.

Table 2: Link prediction using GVAE on the EuCore dataset, trained on the synthetic, tested on real.

Method	Sampling	EuCore	Cora	Wiki	Facebook	SBM
Unif-SAGESS	AUC	0.8463	-	-	0.8601	0.6801
	AP	0.8330	-	-	0.8415	0.6439
RW-SAGESS	AUC	0.8455	0.8952	0.8844	0.8993	0.6921
	AP	0.8340	0.8948	0.8940	0.8862	0.6314
Ego-SAGESS	AUC	0.8463	0.8875	0.8707	0.8889	0.6991
	AP	0.8330	0.8899	0.8856	0.8773	0.6445
NetGAN	AUC	0.8123	0.5099	0.5613	0.8703	0.6978
	AP	0.8156	0.5101	0.5830	0.8499	0.6407
CELL	AUC	0.8566	0.7903	0.8408	0.8714	0.7001
	AP	0.8456	0.8056	0.8637	0.8467	0.6479
DCSBM	AUC	0.6555	0.5096	0.4965	0.4978	0.5378
	AP	0.6383	0.5036	0.5113	0.4976	0.5261

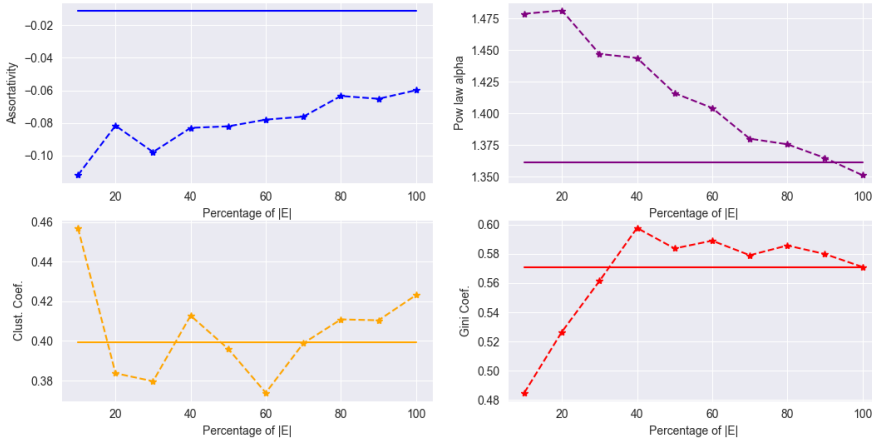


Figure 2: Graphs statistics as a function of percentage of sampled edges from the EuCore data set using the SAGESS-RW model. Note, y -axis is scaled to data, however, the amount of variation of the statistics is relatively small.

Finally, we can also conclude from Figure 2 that our model can generate graphs with a substantially smaller number of edges, while maintaining reasonable approximations of the graph statistics. For example, smaller samples have connectivity and structural properties that are similar to the initial graph even with 40% of its edges. This is not only an interesting feature of our sampling scheme/generator, it is also potentially useful from a computational perspective, when one needs to deploy this methodology on machines with limited resources or shorter time frames, especially as the complexity of graph operations often depends on the number of edges as well as nodes.

6 Discussion and Conclusion

In this paper, we proposed SAGESS, a sampling-based denoising diffusion probabilistic model, based on DIGRESS. It is presented as a framework that is able to build a training data set from a single graph observation, and use it to generate a synthetic graph from a collection of generated synthetic subgraphs. We have shown throughout a variety of experiments that our method outperforms state-of-the-art methods on several real world data sets, not only by considering standard graph statistics, but also by showcasing the adaptability and utility of our framework.

However, this does not come without cost. One limitation of our method is that it still scales poorly in terms of memory with the number of nodes in the initial graph. This is due to the one-hot-encoding of the graph node ids passed in DIGRESS. On the same note, as the number of nodes increases, the number of subgraph samples needed to train our model also increases, adding to the time complexity of DIGRESS. In essence the memory complexity depends on the size of a subsample, resulting in $O(k^2 + kn)$ where k is the size of the subgraph. Additionally, while this method could trivially deal with edge attributes as DIGRESS does, it cannot handle node attributes since they are already taken by the node identification step.

An advantage of our approach is that it can be extended to address other related problems. While this method has been designed to generate graphs from a single sample, it can be extended to data sets with multiple graphs in the case where the nodes are identifiable, by simply designing a sampling scheme which first samples the graph and then samples the subgraph using an appropriate scheme, although we leave this avenue for future work. It will also be left as future work to investigate possible applications on signed networks or even time dependent edges, as one can add time as an edge attribute. This could lead to a more complex generation scheme that could, for instance, bear an auto-regressive module to evaluate time dependence on the newly encountered edges. We would also be eager to extend our framework to other graph generation methods that handle node attributes to generalize the ability of models to train on a single observation.

未来研究方向

Finally, some caution is advised. Inference based on synthetic data particularly in sensitive applications such as health or risk analysis should use a number of different synthetic data generators before reaching a conclusion. Moreover, vigilance is advised to detect malicious applications of synthetic data generation, such as passing off fake data as real.

References

- [1] Réka Albert and Albert-László Barabási. Statistical mechanics of complex networks. *Rev. Mod. Phys.*, 74:47–97, Jan 2002.
- [2] Waqar Ali, Anatol E Wegner, Robert E Gaunt, Charlotte M Deane, and Gesine Reinert. Comparison of large networks with sub-sampling strategies. *Scientific reports*, 6(1):28955, 2016.
- [3] Brett K Beaulieu-Jones, Zhiwei Steven Wu, Chris Williams, Ran Lee, Sanjeev P Bhavnani, James Brian Byrd, and Casey S Greene. Privacy-preserving generative deep neural networks support clinical data sharing. *Circulation: Cardiovascular Quality and Outcomes*, 12(7):e005122, 2019.
- [4] Aleksandar Bojchevski, Oleksandr Shchur, Daniel Zügner, and Stephan Günnemann. Netgan: Generating graphs via random walks. In *International conference on machine learning*, pages 610–619. PMLR, 2018.
- [5] Paul Erdős and Alfréd Rényi. On the evolution of random graphs. *Publ. Math. Inst. Hung. Acad. Sci.*, 5(1):17–60, 1960.
- [6] Faezeh Faez, Yassaman Ommi, Mahdiah Soleymani Baghshah, and Hamid R Rabiee. Deep graph generators: A survey. *IEEE Access*, 9:106675–106702, 2021.
- [7] Xiaojie Guo and Liang Zhao. A systematic survey on deep generative models for graph generation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2022.
- [8] Kilian Konstantin Haefeli, Karolis Martinkus, Nathanaël Perraudin, and Roger Wattenhofer. Diffusion models for graphs benefit from discrete state spaces. *arXiv preprint arXiv:2210.01549*, 2022.
- [9] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 33:6840–6851, 2020.
- [10] Paul W. Holland, Kathryn Blackmond Laskey, and Samuel Leinhardt. Stochastic blockmodels: First steps. *Social Networks*, 5(2):109–137, 1983.
- [11] Brian Karrer and Mark EJ Newman. Stochastic blockmodels and community structure in networks. *Physical review E*, 83(1):016107, 2011.
- [12] Renjie Liao, Yujia Li, Yang Song, Shenlong Wang, Charlie Nash, William L. Hamilton, David Duvenaud, Raquel Urtasun, and Richard Zemel. Efficient graph generation with graph recurrent attention networks. In *NeurIPS*, 2019.

- [13] Julian McAuley and Jure Leskovec. Learning to discover social circles in ego networks. In *Proceedings of the 25th International Conference on Neural Information Processing Systems - Volume 1*, NIPS'12, page 539–547, Red Hook, NY, USA, 2012. Curran Associates Inc.
- [14] Andrew McGregor and Rik Sengupta. Graph Reconstruction from Random Subgraphs. In Mikołaj Bojańczyk, Emanuela Merelli, and David P. Woodruff, editors, *49th International Colloquium on Automata, Languages, and Programming (ICALP 2022)*, volume 229 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 96:1–96:18, Dagstuhl, Germany, 2022. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [15] Luca Rendsburg, Holger Heidrich, and Ulrike Von Luxburg. NetGAN without GAN: From random walks to low-rank approximations. In Hal Daumé III and Aarti Singh, editors, *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 8073–8082. PMLR, 13–18 Jul 2020.
- [16] Garry Robins, Pip Pattison, Yuval Kalish, and Dean Lusher. An introduction to exponential random graph (p^*) models for social networks. *Social Networks*, 29(2):173–191, 2007. Special Section: Advances in Exponential Random Graph (p^*) Models.
- [17] Arne Schneuing, Yuanqi Du, Charles Harris, Arian Jamasb, Ilia Igashov, Weitao Du, Tom Blundell, Pietro Lió, Carla Gomes, Max Welling, et al. Structure-based drug design with equivariant diffusion models. *arXiv preprint arXiv:2210.13695*, 2022.
- [18] Martin Simonovsky and Nikos Komodakis. Graphvae: Towards generation of small graphs using variational autoencoders. In *Artificial Neural Networks and Machine Learning–ICANN 2018: 27th International Conference on Artificial Neural Networks, Rhodes, Greece, October 4-7, 2018, Proceedings, Part I* 27, pages 412–422. Springer, 2018.
- [19] Jascha Sohl-Dickstein, Eric Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. In *International Conference on Machine Learning*, pages 2256–2265. PMLR, 2015.
- [20] Boris van Breugel, Trent Kyono, Jeroen Berrevoets, and Mihaela van der Schaar. Decaf: Generating fair synthetic data using causally-aware generative networks. *Advances in Neural Information Processing Systems*, 34:22221–22233, 2021.
- [21] Clement Vignac, Igor Krawczuk, Antoine Siraudin, Bohan Wang, Volkan Cevher, and Pascal Frossard. Digress: Discrete denoising diffusion for graph generation. *arXiv preprint arXiv:2209.14734*, 2022.
- [22] Sebastien C Wong, Adam Gatt, Victor Stamatescu, and Mark D McDonnell. Understanding data augmentation for classification: when to warp? In *2016 international conference on digital image computing: techniques and applications (DICTA)*, pages 1–6. IEEE, 2016.
- [23] Minkai Xu, Lantao Yu, Yang Song, Chence Shi, Stefano Ermon, and Jian Tang. Geodiff: A geometric diffusion model for molecular conformation generation. *arXiv preprint arXiv:2203.02923*, 2022.
- [24] Renchi Yang, Jieming Shi, Xiaokui Xiao, Yin Yang, Juncheng Liu, and Sourav S Bhowmick. Scaling attributed network embedding to massive graphs. *arXiv preprint arXiv:2009.00826*, 2020.
- [25] Zhilin Yang, William W. Cohen, and Ruslan Salakhutdinov. Revisiting semi-supervised learning with graph embeddings, 2016.
- [26] Hao Yin, Austin R. Benson, Jure Leskovec, and David F. Gleich. Local higher-order graph clustering. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '17, page 555–564, New York, NY, USA, 2017. Association for Computing Machinery.
- [27] Jiaxuan You, Rex Ying, Xiang Ren, William Hamilton, and Jure Leskovec. GraphRNN: Generating realistic graphs with deep auto-regressive models. In *International conference on machine learning*, pages 5708–5717. PMLR, 2018.